

## **ARCO: Technical Implementation & Architectural Decisions**

**System:** Neuro-Symbolic Compliance Verification Engine **Version:** 1.0 (Prototype) **Date:** December 22, 2025

---

### **1. Architectural Decision: Modeling Latent Risk via BFO Dispositions**

**The Problem:** Traditional compliance tags (e.g., "High Risk: True") fail to capture conditional capabilities. A system may have hardware for biometric identification (High Risk) that is currently disabled via software. **The Solution:** We utilize **BFO 2.0 Dispositions** to model capability as a "Realizable Entity" that inheres in the hardware (**Object Aggregate**), independent of its current **Process** realization. **Impact:** This allows the system to flag "Latent Liability" even when the software is configured to "Off."

**Implementation Artifact** ([ARCO\\_core\\_gold\\_fixed\\_v3.ttl](#)):

```
#####
# REALITY-SIDE UNIVERSALS (BFO)
# Decision: Model System as Object Aggregate bearing Dispositions
#####

:System rdf:type owl:Class ;
    rdfs:subClassOf bfo:0000027 ; # Object Aggregate
    rdfs:subClassOf [
        a owl:Restriction ;
        owl:onProperty ro:0000053 ; # bearer of
        owl:someValuesFrom :CapabilityDisposition
    ] .

:BiometricIdentificationCapability rdf:type owl:Class ;
    rdfs:label "Biometric Identification Capability" ;
    rdfs:subClassOf :CapabilityDisposition .
    # Note: This exists whether realized in a process or not.
```

---

### **2. Architectural Decision: Structural Integrity via SHACL**

**The Problem:** OWL Reasoners (Open World Assumption) are poor at validating data completeness. They infer missing data rather than flagging it as an error. **The Solution:** We

inject a **SHACL Validation Layer** (`pyshacl`) prior to the reasoning step. This enforces a "Closed World" constraint on the graph, rejecting any `AssessmentDocumentation` that fails to link a System to its Regulatory Content. **Impact:** Prevents "Garbage In, True Out" scenarios where incomplete data passes audit silently.

**Implementation Artifact** (`assessment_documentation_shape.ttl`):

```
#####
# Shape: Enforce Linkage between System and Regulation
#####

:AssessmentDocumentationShape
  a sh:NodeShape ;
  sh:targetClass :AssessmentDocumentation ;

  sh:property [
    sh:path iao:0000136 ;           # is about
    sh:class :System ;
    sh:minCount 1 ;                # Constraint: Must explicitly link to
System
  ] ;

  sh:property [
    sh:path iao:0000136 ;           # is about
    sh:class :RegulatoryContent ;
    sh:minCount 1 ;                # Constraint: Must explicitly link to
Regulation
  ] .
```

---

### 3. Architectural Decision: Deterministic Audit via SPARQL ASK

**The Problem:** Auditors require binary (Pass/Fail) verdicts. Standard knowledge graph queries (`SELECT`) return lists that require post-processing interpretation, introducing ambiguity. **The Solution:** We utilize **SPARQL ASK** queries which return a boolean `xsd:boolean` value. The query pattern matches the specific absence of compliance or the presence of risk features.

**Impact:** The output is mathematically deterministic and machine-readable (JSON Boolean).

**Implementation Artifact (`ask_assessment_doc_process_wiring.sparql`):**

```
PREFIX : <https://arco.ai/ontology/core#>
PREFIX iao: <http://purl.obolibrary.org/obo/IAO_>

# Objective: Verify if AssessmentDoc connects System to Annex III Condition
ASK WHERE {
    :AssessmentDoc_001 iao:<0000136> :Sentinel_ID_System .
    :AssessmentDoc_001 iao:<0000136> :AnnexIII_Condition_Q1 .
}
# Output: TRUE (Verification Successful)
```

---

#### 4. Execution Pipeline (Neuro-Symbolic Bridge)

**Orchestration:** The Python runtime (`rdflib` + `pyshacl`) manages the transition from probabilistic extraction (LLM-generated instances) to deterministic validation.

**Pipeline Code (`run_checks.py`):**

```
Python
def run_checks(data_graph: Graph):
    # Step 1: Structural Validation (SHACL)
    conforms, _, report_text = validate(
        data_graph=data_graph,
        shacl_graph=Graph().parse(SHAPES),
        inference="none"
    )
    if not conforms:
        raise ValidationError(f"Graph Structure Invalid: {report_text}")

    # Step 2: Compliance Reasoning (SPARQL)
    # Only executes if Structure is Valid
    run_sparql_audit(data_graph)
```